

Machine Learning Fundamentals

Week 1

Task 1 – Exploratory Data Analysis (EDA)

Dataset: Titanic - Machine Learning from Disaster

Prepared by: Abdul Raheem

Track: Neurofive Machine Learning Track

Objective

The objective of this notebook is to explore the Titanic dataset, understand its structure, identify missing values, distinguish numerical and categorical features, and summarize the dataset before applying any machine learning models.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Importing the Dataset

You can use this instruction `df = pd.read_csv("train.csv")`, only if Python can find `train.csv` in the current working directory (the folder from which Jupyter is running). Means `Week1_Titanic_EDA.ipynb` and `train.csv` are in the same folder.

But if the dataset and the notebook are in different folders then use the following approach.

```
df = pd.read_csv(r"D:\Machine Learning Fundamentals Course\Week 1\
titanic\train.csv")
```

```
df.head()
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

SibSp	\	Name	Sex	Age
0		Braund, Mr. Owen Harris	male	22.0

```

1
1 Cumings, Mrs. John Bradley (Florence Briggs Th... female 38.0
1
2 Heikkinen, Miss. Laina female 26.0
0
3 Futrelle, Mrs. Jacques Heath (Lily May Peel) female 35.0
1
4 Allen, Mr. William Henry male 35.0
0

```

```

    Parch      Ticket    Fare Cabin Embarked
0      0    A/5 21171    7.2500   NaN        S
1      0    PC 17599   71.2833   C85        C
2      0 STON/O2. 3101282    7.9250   NaN        S
3      0    113803   53.1000  C123        S
4      0    373450    8.0500   NaN        S

```

```
df.info()
```

```

<class 'pandas.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   str
4   Sex             891 non-null   str
5   Age             714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   str
9   Fare            891 non-null   float64
10  Cabin           204 non-null   str
11  Embarked        889 non-null   str
dtypes: float64(2), int64(5), str(5)
memory usage: 83.7 KB

```

```
df.describe()
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	20.125000	0.000000	
50%	446.000000	0.000000	3.000000	28.000000	0.000000	
75%	668.500000	1.000000	3.000000	38.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
print(df.shape)
```

```
(891, 12)
```

```
df.isnull().sum()
```

```

PassengerId      0
Survived          0
Pclass            0
Name              0
Sex               0
Age              177
SibSp             0
Parch             0
Ticket            0
Fare              0
Cabin            687
Embarked          2
dtype: int64

```

```
df.select_dtypes(include=np.number).columns
```

```
Index(['PassengerId', 'Survived', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare'], dtype='str')
```

```
df.select_dtypes(include=['object', 'str']).columns
```

```
Index(['Name', 'Sex', 'Ticket', 'Cabin', 'Embarked'], dtype='str')
```

Data Story

The Titanic dataset contains information about 891 passengers and 12 different features. The dataset includes both numerical and categorical variables such as age, fare, gender, passenger class, and embarkation point. Some columns, including Age, Cabin, and Embarked, contain missing values that will require preprocessing before training machine learning models. Overall, the dataset appears well-structured and provides sufficient information for predicting passenger survival.

Task 2 - Data Cleaning & Visualization

After understanding the dataset through Exploratory Data Analysis (EDA), the next step is to prepare the data for Machine Learning. This includes handling missing values, identifying potential outliers, and visualizing the dataset to better understand relationships among features.

The objective of this task is to clean the Titanic dataset, create meaningful visualizations, and summarize the insights obtained from the analysis.

Handling Missing Values

Before building any Machine Learning model, missing values must be addressed because they can reduce model performance and may even prevent certain algorithms from working correctly. The assignment requires using `fillna()` or `dropna()` and justifying the choice.

First, we inspect the dataset to identify which columns contain missing values.

```
df.isnull().sum()
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age            177
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin          687
Embarked         2
dtype: int64

df["Age"] = df["Age"].fillna(df["Age"].median())
df["Embarked"] = df["Embarked"].fillna(df["Embarked"].mode()[0])
df["Cabin"] = df["Cabin"].fillna("Unknown")
```

Why was `fillna()` used?

Instead of removing rows using `dropna()`, missing values were filled because dropping rows would reduce the size of the dataset and could result in the loss of useful information.

The median was used for the **Age** column because it is less affected by extreme values than the mean. The mode was used for **Embarked** because it is a categorical feature. Missing values in **Cabin** were replaced with "Unknown" since most cabin entries are unavailable. Now Verify the Missing Values again.


```
df.isnull().sum()
```

```
PassengerId    0  
Survived        0  
Pclass         0  
Name           0  
Sex            0  
Age           0  
SibSp         0  
Parch         0  
Ticket        0  
Fare          0  
Cabin         0  
Embarked       0  
dtype: int64
```

Histogram

A histogram illustrates the distribution of passenger ages.


```
plt.figure(figsize=(8,5))  
sns.histplot(df["Age"], bins=30, kde=True)  
plt.title("Age Distribution")  
plt.xlabel("Age")  
plt.ylabel("Frequency")  
plt.show()
```




Outlier Detection

A boxplot is used to identify potential outliers in passenger fares.

```
plt.figure(figsize=(8,5))
sns.boxplot(x=df["Fare"])
plt.title("Fare Boxplot")
plt.show()
```




```
plt.figure(figsize=(7,5))
sns.countplot(x="Survived", data=df)
plt.title("Survival Count")
plt.show()
```



```
df_encoded = df.copy()
df_encoded["Sex"] = df_encoded["Sex"].map({"male":0,"female":1})
df_encoded["Embarked"] =
df_encoded["Embarked"].map({"S":0,"C":1,"Q":2})

plt.figure(figsize=(10,8))
sns.heatmap(df_encoded.corr(numeric_only=True), annot=True,
cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()
```

Which feature most affects survival?

Based on the visualizations and correlation analysis, **Sex** appears to have the strongest influence on passenger survival. Female passengers had a noticeably higher survival rate than male passengers. This aligns with the historical "women and children first" evacuation policy followed during the Titanic disaster.

Passenger class (Pclass) also shows a significant relationship with survival, indicating that first-class passengers generally had a higher chance of surviving than those in lower classes.

Conclusion

The Titanic dataset was successfully cleaned by handling missing values using appropriate techniques for numerical and categorical features. Visualizations provided valuable insights into the distribution of data, potential outliers, and relationships among variables.

The analysis indicates that passenger gender and class were among the most influential factors affecting survival. These findings provide a solid foundation for the predictive modeling tasks in the upcoming weeks.

